

Aadarsh Upadhyay

Software Engineer / Architect

Legal name: Aadarsh Raj

India · remote-first · open to relocate · outside-India visa sponsorship welcome

capad.fyi · github.com/capad-xyz

hi@capad.fyi · x.com/aadarsh_io

[linkedin.com/in/aadarshupadhyay](https://www.linkedin.com/in/aadarshupadhyay)

React-first full-stack engineer. Home stack: CSS, React, TypeScript, Next.js, Electron, Node, deploy. Product depth from Appson on **ComplyV** (formerly Compliance Sarathi) and **Wordibly**. Public proof: **searchts** on PyPI through **0.13.0** (I push, review, and release). I also own [Grove](#), **Doooper**, **GlyphMaps**, and **capad.fyi**; Hare is not a SaaS yet; shipped **Hare Bot**. I use AI assistants to ship and review hard. I am finishing BCA (Honours) at MSU Baroda (expected 2028). Looking for remote or India roles that hire for ownership of shipped systems. Open to relocate; outside-India roles with visa sponsorship welcome.

EXPERIENCE

Software Engineer and Architect, Appson Technologies (ComplyV, formerly Compliance Sarathi) Jan 2026 - 31 Jul 2026

Indian **corporate-governance SaaS**: MCA/ROC company data, statutory deadlines and penalties, resolution and minutes drafting, AI chat (Vee) with governed writes. Dual UI (legacy CRA + /new-ui). **Primary engineer and architect** on a small team (~553/647 commits (~84%)). Not sole developer. Last commit and written handover 31 Jul 2026. Stack: **React 18** (CRA), React Router, Tailwind, Bootstrap, **Node/Express**, **MongoDB/Mongoose** (34 models), **JWT/RBAC**, WebSocket, OpenAI/Gemini/Anthropic, AWS S3, Twilio, Brevo/Resend, Word/Excel/PDF.

- Owned the **React and Node product surface**: auth, multi-tenant company access, deadline math, minutes/docs drafting, Vee agentic chat. Commit share ~553/647 (~84%).
- Built **propose-then-confirm writes**. The model never mutates data. Proposals show as confirmation cards. A separate authenticated endpoint re-checks auth, company access, and business rules with **idempotency keys** and **audit logs**.
- Built a **three-tier truth wall**. Due dates and penalties come only from the **deadline engine**. Statutory drafts reuse the product drafting pipeline. Eval cases bait **number leakage** on conceptual answers.
- Implemented a **SHA-256 append-only audit log** with immutability hooks, deterministic hashing, a write mutex, and an **admin integrity verifier**. Tamper-resistant, not tamper-proof.
- Multi-tenancy** without a naive tenantId. Company access is an **ownership-and-grant graph** over CIN. LLM tools get a CIN string that resolves inside **assertCompanyAccess**.
- Maintained the **AGM/deadline engine** (~1,400 LOC of pure computation) for multiple MCA forms with entity-aware deadlines and per-day penalties. Draft linter is code, not model self-grading. Fixed **LLP fee slabs** and a **DIR-3 KYC triennial rule** that would have been wrong from 2027.
- Registered **34 Mongoose models** across identity, tenancy, documents, AI runtime, governance, and commerce. Backend tests: ~40 suites / ~2,852 tests (technical sheet).
- Found **auth gaps** while designing agent tools (**deny-by-default** draft visibility and writes). Made **idempotency a unique index in the database**, not an in-memory lock.
- Wrote the 31 Jul 2026 handover: architecture, known issues, secret rotation.

Junior Software Engineer, Appson Technologies (Wordibly)

Jun 2025 - early 2026

US **transcription, translation, and captioning platform** (human, hybrid, AI tiers). Strongest personal ownership: **tops-transcript-editor (Vite/React)**. Broader platform (tops-api, tops-new-ui, tops-old-ui) is large-team CodeCommit with thinner personal share. Stack exposure: React, Node, Python, MongoDB, AWS; Zoho, QuickBooks, Zapier, Otter.

- Owned **tops-transcript-editor** (~107-113 commits): **Vite/React transcript editor** with find/replace and related editing flows for production transcription work.
- Contributed on the wider **CodeCommit platform** (tops-api, tops-new-ui, tops-old-ui) with a thinner personal share: **order/upload surfaces**, **workforce dashboards**, and **editor workflows** for PMs, transcribers, and QA.
- Daily Scrum with Appson and Wordibly. **Sprint delivery** with **no visual regressions** on production-bound UI.

AI Training Engineer, Turing (via Appson)

~2 weeks, 2025

Meta Llama training data and evaluation.

OPEN SOURCE AND PRODUCTS

searchts (Python, MIT, PyPI v0.13.0 · github.com/capad-xyz/searchts)

My product. Keyless web unlocker plus read/search/transcribe/grab for AI agents (**CLI + MCP**). I own the unlocker, **search fusion**, sanitizer, MCP tools, share-link extractors, and **releases** (push, review, ship through 0.13.0). The local MCP/connector for searchts is built; the remote connector will be built soon. Escalates browser-fingerprinted fetch → JS-render relay → stealth browser, then Markdown extraction. Defaults refuse paid residential proxies and keyed commercial unlockers so personal-volume use stays free. Fail-loud on login walls and CAPTCHAs. Agent-facing docs/skill pack (AGENTS.md, RUNBOOK, CLAUDE.md, SKILL.md) so assistants use searchts the intended way, not as a raw fetch wrapper.

Hare (not a SaaS yet · GitHub App **searchts-hare**[bot] · @hare · R1c · dogfooded on searchts)

Not a SaaS yet: a **review-and-report GitHub App** I dogfood on searchts. I design the full **PR review surface**: structured summary, severity findings, inline bubbles, and replies, grounded in **CI and diff evidence**, non-blocking feedback without a SaaS lock-in. Trigger: PR open, push, or /hare → gather → **one model JSON** → posted review (R1c). Writer model never Hare model. **Review/report only**; never blocking APPROVE/REQUEST_CHANGES. Not a SaaS yet; no CodeRabbit parity claims.

Hare Bot (Cursor Grok Bot · PR review/report)
Maintainers need **structured non-blocking PR review** without SaaS lock-in; agent desks need review loops they can own. I design those loops as a Cursor Grok Bot: reads **PR body/diff/commits/CI as evidence**, can checkout and run tests, then posts a **COMMENT review** with summary, severity findings, inline bubbles, replies, and Models attribution: more of the review loop than a single dump. **Review/report only**. Hare Bot does not fix, push, or merge unless asked; never blocking APPROVE/REQUEST_CHANGES. Chat-side template; Hare App is a separate surface.

capad.fyi (my site · Next.js, React, Sanity, Cloudflare Workers / OpenNext · [capad.fyi](#))
My site and product desk: **liquid-glass UI** (CSS backdrop-filter + SVG refraction), **Sanity-backed projects / timeline / testimonials / resume** with HMAC webhook revalidation into KV-backed ISR, multi-host Worker for capad.fyi and glyphmaps.capad.fyi, adaptive perf-lite path, **GitHub Actions deploy**.

MSU Halls Register (design prototype · [msu.capad.fyi](#))
Institutional alumni-register prototype for the Halls of Residence at The Maharaja Sayajirao University of Baroda (**16 halls**). **Eight-page static design system** revised after Office of the Chief Warden feedback. Live prototype; engagement paused, not under active contract.

GlyphMaps (Kotlin / Android, v1.0.0 · [github.com/capad-xyz/GlyphMaps](#))
Google Maps Live Updates on the **Nothing Glyph Matrix**, built for Nothing Phone (4a) Pro; works on Phone 3 (needs work). **Next-turn arrows** on the phone back, **no Maps API key**. Skipped the throttled Glyph Toy redraw path. Drives the matrix while navigating, then releases it so other Glyph toys return.

Dooper (repo beep-beep-oss · Tauri 2 + React/TypeScript + Matrix · Alpha Phase 1)
AGPL self-hostable universal inbox: Synapse, mautrix bridges, Tauri 2 + React. Messaging sync without paywalled delay. **Phase 1 verified on real bridged WhatsApp** (login, inbox, history, **optimistic send**, **session persistence**). Product name Dooper; repo is beep-beep-oss.

Grove (Electron + React + TypeScript, GPL-3.0 · alpha · [github.com/capad-xyz/grove](#))
Rebuilt on Electron, React, and TypeScript with a **headless Node git engine**. **Worktree-aware git review** beside AI editors: lane-drawn commit graph, live refresh, **find-in-diff**, **local-CLI commit messages**. No paid in-house model lock-in.

Also: burncard (npx burncard, local Claude Code / Codex usage telemetry).

SELECTED OPEN-SOURCE CONTRIBUTIONS

wmux (contributor · [amirlehman/wmux](#) · not my project)
I dogfood wmx in daily agent work. I file issues when I hit bugs; I land PRs when I have a fix. As of 2026-10-02: **8 closed issues**, **4 merged PRs** ([#135](#), [#138](#), [#153](#), [#258](#)) on **diff-pane freezes**, **CLI timeouts**, and agent-browser install discovery. Best proof: [#135](#) — root-caused a **main-process freeze** when a diff pane watched a non-git cwd; landed **bounded snapshot walks**.

EDUCATION

Bachelor of Computer Applications (Honours), The Maharaja Sayajirao University of Baroda, Faculty of Science, Dept. of Computer Applications	Expected 2028
Product Space PM Fellowship, Top Fellow	2025 - 2026

SKILLS

Tech stack	React, React Three Fiber, CSS, design systems, Tailwind, Bootstrap, Framer Motion, React Router, Next.js (App Router), Vite, Electron. Layout through production deploy.
Design / UI-UX	UI/UX, Figma, design systems, product UI from layout through ship. AI-assisted design drafts with Claude Design and Stitch. Proof: ComplyV /new-ui , Wordibly transcript editor, capad.fyi liquid-glass, MSU Halls register prototype, GlyphMaps matrix UI.
Full-stack	TypeScript, JavaScript, Node.js, Express, MongoDB/Mongoose, REST/JWT, Python (CLI), Git, GitHub Actions, MCP, Cloudflare Workers / OpenNext, Sanity
AI / tools	Grok / Grok Bots (incl. Hare Bot), Claude; daily drivers Oh My Pi (OMP) and GrokCLI; MCP, T3 Connect, Cowork, Custom Connectors, WebMCP; local searchts MCP/connector built (remote soon). Custom skills / AGENTS.md / RUNBOOK; writer ≠ Hare reviewer; PR text/diffs/CI as evidence , not instructions.
AI shipping	I own push and review on searchts and other repos: product decisions , release notes , merge-ready diffs , documented refusals of paid unlocker defaults
Agent skills	Custom agent skills, runbooks, AGENTS.md, and RUNBOOK so multi-agent desks follow product constraints instead of improvising
Systems	Agent tool design, authz, multi-tenancy, audit logs, eval harnesses, prompt-vs-code safety boundaries
Ship / DevOps	GitHub Actions (CI, release, Hare review on searchts), Cloudflare Workers / OpenNext (capad.fyi), Docker, Tailscale, AWS S3. Comfortable operating what I ship; not applying as a dedicated DevOps engineer yet.
Also	Tauri, Sanity CMS, Twilio, Brevo/Resend; Rust/Kotlin with AI assist (Dooper, GlyphMaps). Not targeting Rust/Kotlin/Go-primary roles.